

# Homework 5

Yixin Wang

## Contents

Elastic Net Tuning . . . . .	1
------------------------------	---

## Elastic Net Tuning

The goal of this assignment is to build a statistical learning model that can predict the **primary type** of a Pokémon based on its generation, legendary status, and six battle statistics.

Read in the file and familiarize yourself with the variables using `pokemon_codebook.txt`.

```
library(tidymodels)
library(tidyverse)
library(discrim)
library(poissonreg)
library(ISLR)
library(ISLR2)
library(corr)
library(klaR)
library(glmnet)
tidymodels_prefer()
set.seed(3435)
```

## Exercise 1

Install and load the `janitor` package. Use its `clean_names()` function on the Pokémon data, and save the results to work with for the rest of the assignment. What happened to the data? Why do you think `clean_names()` is useful?

```
library(janitor)
```

```
pokemon <- read_csv(file = "Pokemon.csv")
pokemon <- clean_names(pokemon)
pokemon <- as_tibble(pokemon)
pokemon
```

```
## # A tibble: 800 x 13
##   number name      type_1 type_2 total    hp attack defense sp_atk sp_def speed
##   <dbl> <chr>    <chr> <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1     1 1 Bulbasaur Grass Poison 318   45   49   49    65    65   45
## 2     2 2 Ivysaur   Grass Poison 405   60   62   63    80    80   60
```

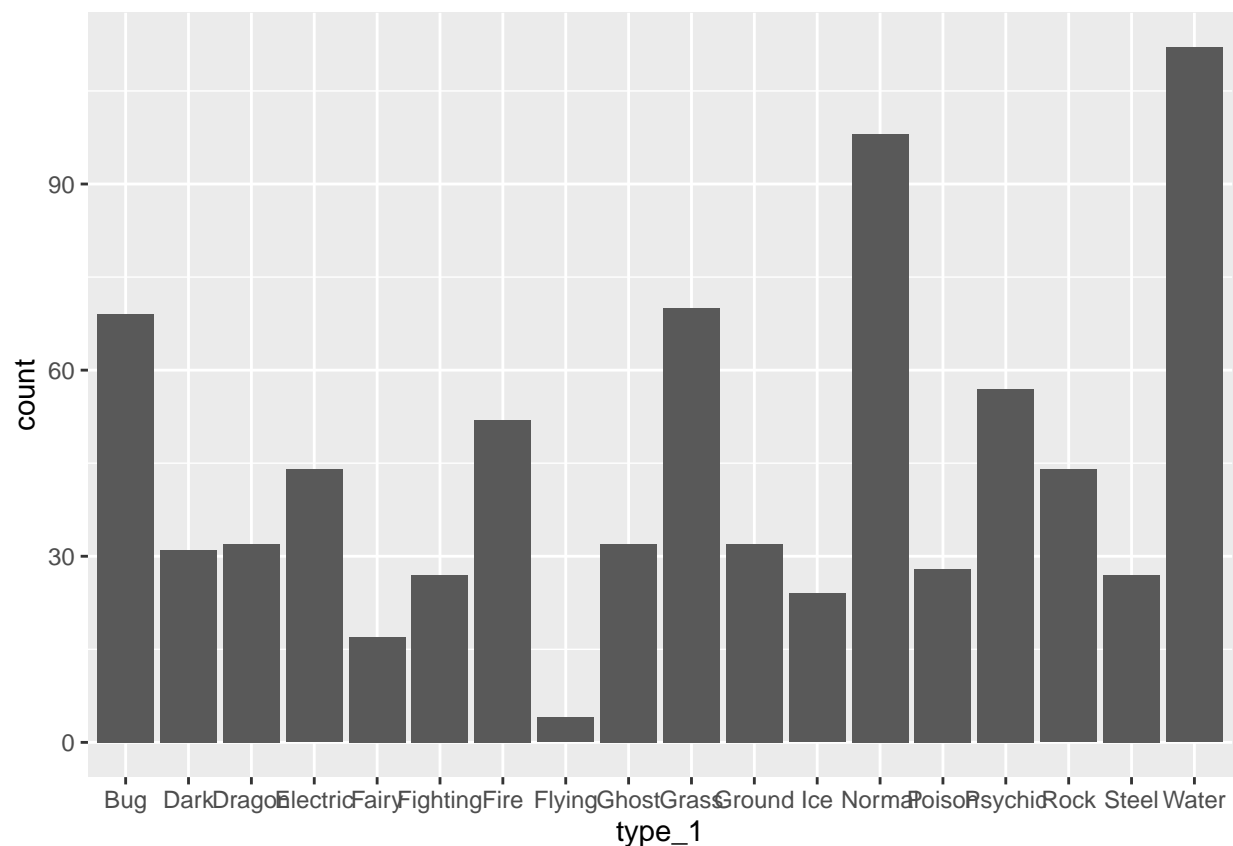
```
## 3      3 Venusaur  Grass  Poison  525    80    82    83    100    100    80
## 4      3 Venusaur~ Grass  Poison  625    80   100   123   122   120    80
## 5      4 Charmand~ Fire   <NA>   309    39    52    43    60    50    65
## 6      5 Charmele~ Fire   <NA>   405    58    64    58    80    65    80
## 7      6 Charizard Fire   Flying  534    78    84    78   109    85   100
## 8      6 Charizar~ Fire   Dragon  634    78   130   111   130    85   100
## 9      6 Charizar~ Fire   Flying  634    78   104    78   159   115   100
## 10     7 Squirtle  Water  <NA>    314    44    48    65    50    64    43
## # ... with 790 more rows, and 2 more variables: generation <dbl>,
## #   legendary <lgl>
```

The data handles problematic variable names. The `clean_names()` function cleans the variable names and return to data frame.

## Exercise 2

Using the entire data set, create a bar chart of the outcome variable, `type_1`.

```
library(ggplot2)
ggplot(pokemon, aes(x = type_1)) + geom_bar()
```



How many classes of the outcome are there? Are there any Pokémon types with very few Pokémon? If so, which ones?

There are 18 classes of the outcome. The flying and fairy types are have few Pokemons.

For this assignment, we'll handle the rarer classes by simply filtering them out. Filter the entire data set to contain only Pokémon whose `type_1` is Bug, Fire, Grass, Normal, Water, or Psychic.

After filtering, convert `type_1` and `legendary` to factors.

```
pokemon <- pokemon[pokemon$type_1 %in% c("Bug", "Fire", "Grass", "Normal", "Water", "Psychic"), ]
```

```
pokemon$type_1 <- as.factor(pokemon$type_1)
pokemon$legendary <- as.factor(pokemon$legendary)
```

### Exercise 3

Perform an initial split of the data. Stratify by the outcome variable. You can choose a proportion to use. Verify that your training and test sets have the desired number of observations.

Next, use *v*-fold cross-validation on the training set. Use 5 folds. Stratify the folds by `type_1` as well. *Hint: Look for a `strata` argument.* Why might stratifying the folds be useful?

```
pokemon_split <- initial_split(pokemon, prop = 0.7, strata = type_1)
pokemon_train <- training(pokemon_split)
pokemon_test <- testing(pokemon_split)
```

```
dim(pokemon_train)
```

```
## [1] 318 13
```

```
dim(pokemon_test)
```

```
## [1] 140 13
```

```
pokemon_fold <- vfold_cv(pokemon_train, v = 5, strata = type_1)
```

Since the outcome variable is imbalanced, there are less observations for flying and fairy than other types.

### Exercise 4

Set up a recipe to predict `type_1` with `legendary`, `generation`, `sp_atk`, `attack`, `speed`, `defense`, `hp`, and `sp_def`.

- Dummy-code `legendary` and `generation`;
- Center and scale all predictors.

```
pokemon_recipe <- recipe(type_1 ~ legendary + generation + sp_atk + attack + speed + defense + hp + sp_
  step_dummy(all_nominal_predictors()) %>%
  step_scale(all_predictors()) %>%
  step_center(all_predictors())
```

## Exercise 5

We'll be fitting and tuning an elastic net, tuning `penalty` and `mixture` (use `multinom_reg` with the `glmnet` engine).

Set up this model and workflow. Create a regular grid for `penalty` and `mixture` with 10 levels each; `mixture` should range from 0 to 1. For this assignment, we'll let `penalty` range from -5 to 5 (it's log-scaled).

```
pokemon_spec <- multinom_reg(penalty = tune(), mixture = tune()) %>%  
  set_engine("glmnet") %>%  
  set_mode("classification")  
  
pokemon_grid <- grid_regular(penalty(c(-5, 5)), mixture(c(0,1)), levels = 10)  
  
pokemon_workflow <- workflow() %>%  
  add_recipe(pokemon_recipe) %>%  
  add_model(pokemon_spec)
```

How many total models will you be fitting when you fit these models to your folded data?

There are 10 models will be fitted to 5 folded data.

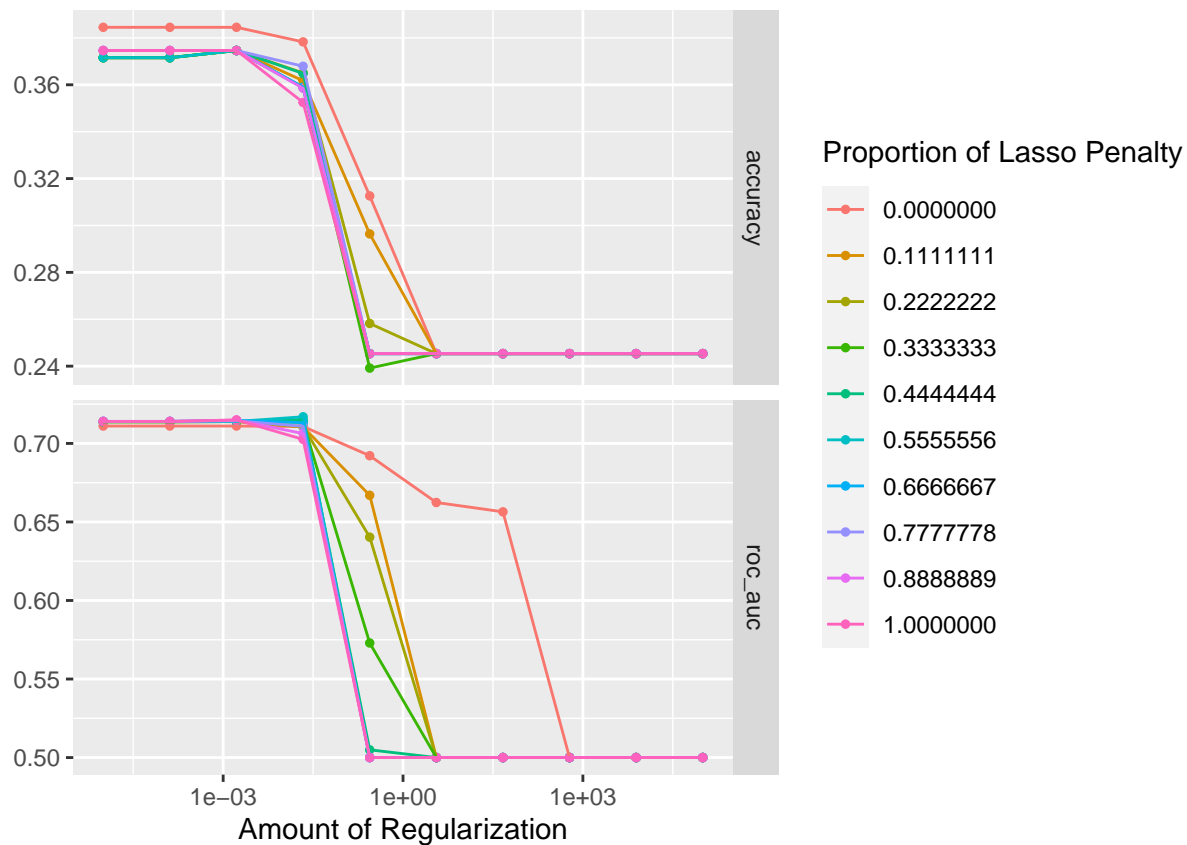
## Exercise 6

Fit the models to your folded data using `tune_grid()`.

```
tune_res <- tune_grid(pokemon_workflow, resamples = pokemon_fold, grid = pokemon_grid)
```

Use `autoplot()` on the results. What do you notice? Do larger or smaller values of `penalty` and `mixture` produce better accuracy and ROC AUC?

```
autoplot(tune_res)
```



The smaller the penalty lead to the better accuracy and ROC AUC. Similarly, the smaller mixture leads to the better accuracy and ROC AUC.

### Exercise 7

Use `select_best()` to choose the model that has the optimal `roc_auc`. Then use `finalize_workflow()`, `fit()`, and `augment()` to fit the model to the training set and evaluate its performance on the testing set.

```
pokemon_penalty <- select_best(tune_res, metric = "roc_auc")
```

```
final = pokemon_workflow %>%
  finalize_workflow(pokemon_penalty) %>%
  fit(pokemon_train) %>%
  augment(pokemon_test)

accuracy(final, truth = type_1, estimate = .pred_class)
```

```
## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>      <dbl>
## 1 accuracy multiclass    0.364
```

### Exercise 8

Calculate the overall ROC AUC on the testing set.

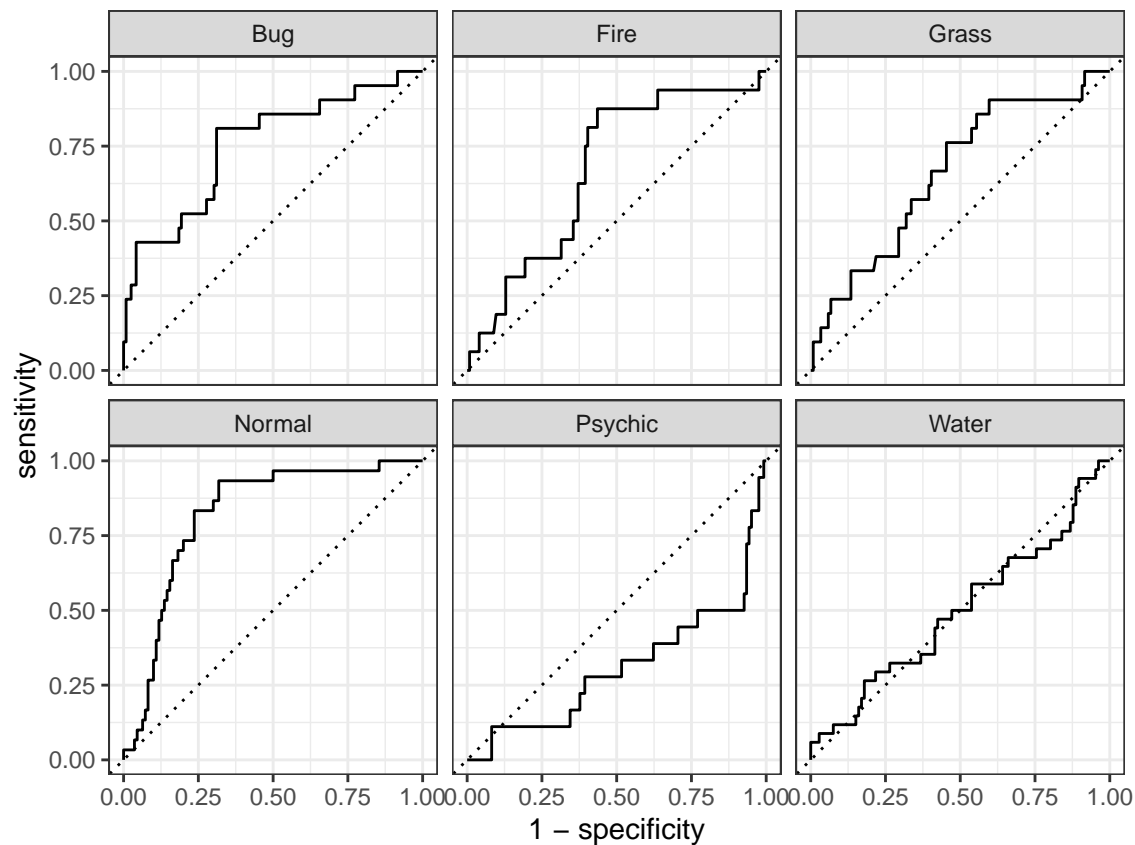
```
roc_auc(final, truth = type_1, .pred_Bug, .pred_Fire, .pred_Grass, .pred_Normal, .pred_Water, .pred_Psy
```

```
## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 roc_auc hand_till     0.615
```

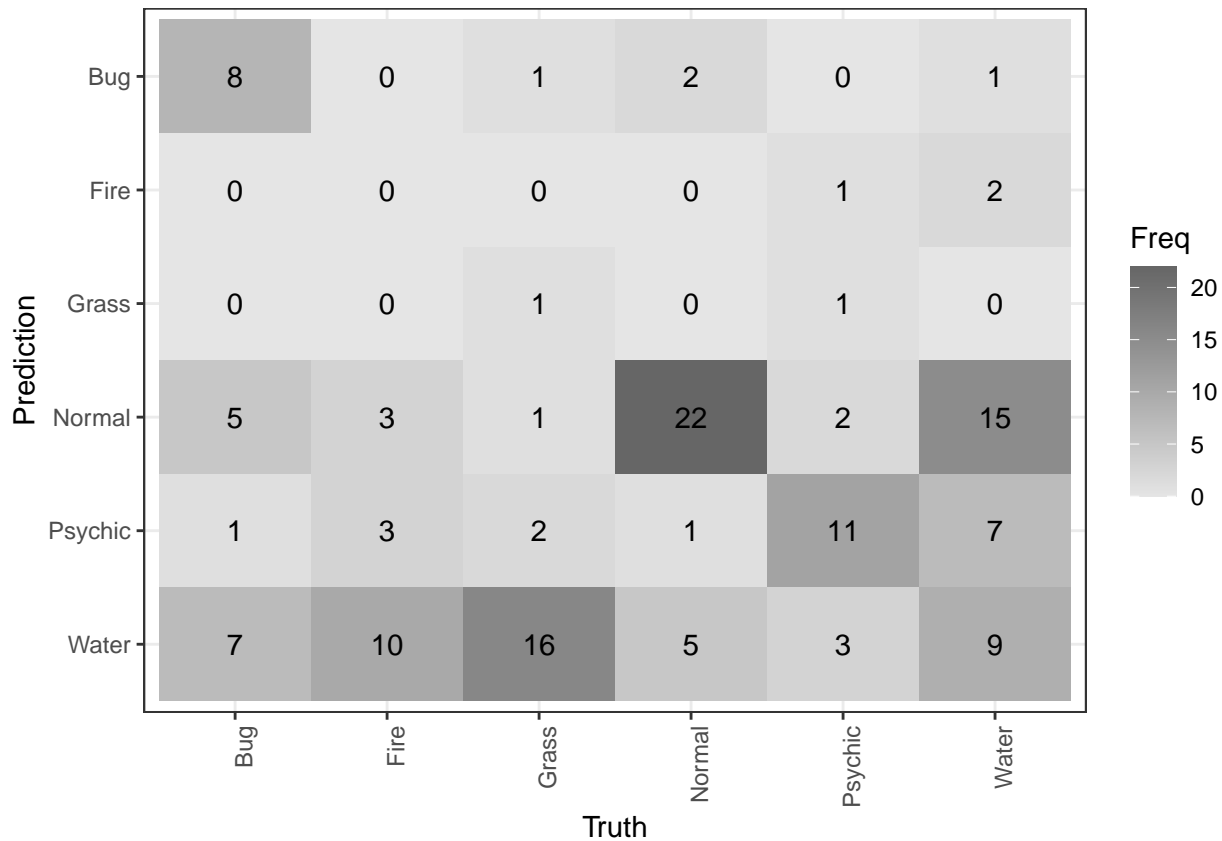
Then create plots of the different ROC curves, one per level of the outcome. Also make a heat map of the confusion matrix.

What do you notice? How did your model do? Which Pokemon types is the model best at predicting, and which is it worst at? Do you have any ideas why this might be?

```
autoplot(roc_curve(final, truth = type_1, .pred_Bug, .pred_Fire, .pred_Grass, .pred_Normal, .pred_Water
```



```
conf_mat(final, truth = type_1, estimate = .pred_class) %>%
  autoplot(type = "heatmap") + theme_bw() + theme(axis.text.x = element_text(angle = 90, hjust = 1))
```



The roc\_auc is approximately 0.6072. The accuracy is approximately 0.45. I notice that the model is not good at predicting the type of Pokemon. However, the normal type is the model best at predicting. The psychic type is the worst type to predict in the graph. I think the reason is probably that the psychic is quite different than the other Pokemons.